



# STATS / DATA SCI 315

## Lecture 03

Regression  
Loss functions and gradient descent

---

# Loss functions and gradient descent



## Loss function

- We need a way to measure how good a prediction  $\hat{y}^{(i)}$  is when the true label is  $y^{(i)}$
- Most popular regression loss is the *squared error*:

$$l^{(i)}(\mathbf{w}, b) = \frac{1}{2} (\hat{y}^{(i)} - y^{(i)})^2$$

- The  $\frac{1}{2}$  above is just for convenience
- Quality of model on entire dataset is assessed by:

$$L(\mathbf{w}, b) = \frac{1}{n} \sum_{i=1}^n l^{(i)}(\mathbf{w}, b) = \frac{1}{n} \sum_{i=1}^n \frac{1}{2} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)})^2.$$

## Visualizing the fit with 1-dimensional $x$




Fig. 3.1.1 Fit data with a linear model.



## Minimizing the loss

- Bias can be absorbed into weights by using a “dummy” feature which is always 1:

$$\mathbf{w}^\top \mathbf{x} + b = (\mathbf{w}, b)^\top (\mathbf{x}, 1)$$

- Best choice for  $\mathbf{w}$  is given by:

$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} L(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}} \frac{1}{2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2$$

- Note that, in the last step, we removed  $n$  and used the Euclidean norm
- This *optimization problem* turns out to have a closed form solution



## Taking derivatives

- How do we take the derivative of  $L(\mathbf{w}) = \frac{1}{2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2$  w.r.t.  $\mathbf{w}$ ?
- View  $L(\mathbf{w})$  as  $F(G(\mathbf{w}))$  and use *chain rule*:

$$L'(\mathbf{w}) = G'(\mathbf{w}) \times F'(G(\mathbf{w}))$$

- $G(\mathbf{u}) = \mathbf{y} - \mathbf{X}\mathbf{u}$  and  $F(\mathbf{u}) = \frac{1}{2} \|\mathbf{u}\|^2$
- $G'(\mathbf{u}) = -\mathbf{X}^T$  and  $F'(\mathbf{u}) = \mathbf{u}$
- Therefore, the derivative w.r.t.  $\mathbf{w}$  is:

$$(-\mathbf{X}^T) \times (\mathbf{y} - \mathbf{X}\mathbf{w}) = -\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \mathbf{w}$$



## Set gradient to zero

- Note that the derivative w.r.t.  $\mathbf{w}$  has  $d$  components
- It is often called a *gradient*

$$\nabla L(\mathbf{w}) = -\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \mathbf{w}$$

- Setting it to zero gives the closed-form solution

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- We're assuming here that  $\mathbf{X}^T \mathbf{X}$  is invertible



## Closed-form solutions are rare

- We got lucky in our simple setting (linear regression with squared loss)
- Usually we're not so lucky
- Even simple changes can destroy this property (like using absolute error)
- If we insisted on closed form solutions, almost *all* of DL will be excluded
- Key technique for incrementally lowering the loss function: *gradient descent*
- Iteratively reduces the loss by updating the parameters in the direction of the *negative gradient*





# Gradient Descent

- For any objective function  $J(\mathbf{w})$ , GD update takes the form:  
$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla J(\mathbf{w})$$
- The gradient gives you the direction of fastest local *increase* in  $J$
- Since we're looking to minimize  $J$ , we move in the direction of *negative* gradient
- The step size (aka learning rate)  $\eta$  controls how much we move



## Gradient Descent on the Loss

- Takes the form:  
 $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla L(\mathbf{w})$
- Here, the gradient has the form  
 $\nabla L(\mathbf{w}) = 1/n \sum_i \frac{1}{2} \nabla (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2 = 1/n \sum_i \mathbf{x}^{(i)} (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})$
- So GD update becomes  
 $\mathbf{w} \leftarrow \mathbf{w} - \eta/n \sum_i \mathbf{x}^{(i)} (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})$
- Requires one full pass through the entire data set



## Minibatch Stochastic Gradient Descent

- In each iteration, we first randomly sample a minibatch  $\mathbf{B}$  consisting of a fixed number of training examples
- Then we update
$$\mathbf{w} \leftarrow \mathbf{w} - \eta/|\mathbf{B}| \sum_{i \in \mathbf{B}} \mathbf{x}^{(i)} (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})$$
- Note that the set  $\mathbf{B}$  is random and changes from iteration to iteration
- $\eta$  and  $|\mathbf{B}|$  (the batch size) are *hyperparameters*: they're kept fixed during training
- However, an outer loop might optimize them by tracking performance on a *validation set*



## Updates with bias kept separately

- Update  $\mathbf{w}$ :  
$$\mathbf{w} \leftarrow \mathbf{w} - \eta/|\mathbf{B}| \sum_{i \in \mathbf{B}} \mathbf{x}^{(i)} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)})$$
- Update bias:  
$$b \leftarrow b - \eta/|\mathbf{B}| \sum_{i \in \mathbf{B}} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)})$$



## Making predictions with learned model

- Train for some predetermined number of iterations (or until some other stopping criteria are met)
- Record the estimated model parameters, denoted  $\hat{\mathbf{w}}, \hat{b}$
- Given a new house with area  $x_1$  and age  $x_2$ , we predict its price as  $\hat{\mathbf{w}}^\top \mathbf{x} + \hat{b}$  where  $\mathbf{x} = (x_1, x_2)$
- Estimating targets given features is commonly called *prediction*
- It's also (misleadingly) called *inference* in deep learning jargon