

Notebook credit: based on the F. Chollet's original notebook [here](#).

✓ Fundamentals of machine learning

✓ Generalization: The goal of machine learning

- we have looked at predicting movie reviews and house-price regression
- we split the data into a training set, a validation set, and a test set
- we should not evaluate the models on the same data they were trained
- after just a few epochs, performance on never-before-seen data started diverging from performance on the training data
- training data performance always improves as training progresses
- at some point, the models started to **overfit**

The fundamental issue in ML is the tension between optimization and generalization

- optimization: process of adjusting a model to get the best performance possible on the training data
- generalization: how well the trained model performs on data it has never seen before
- key challenge: our *goal* is good generalization we can't directly control generalization
- can only fit the model to its training data
- do that too well: overfitting kicks in and generalization suffers

Underfitting and overfitting

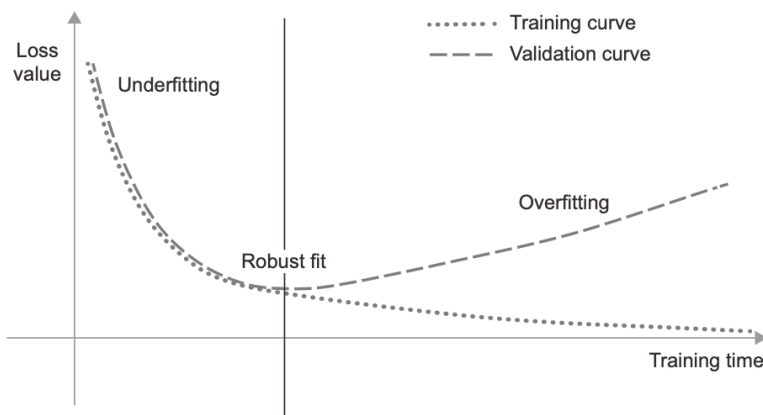


Figure 5.1 Canonical overfitting behavior

✓ 什么时候容易发生 Overfitting

Overfitting is particularly likely to occur when your data:

- is noisy

- involves uncertainty
- includes rare features

Adding white-noise channels or all-zeros channels to MNIST

- we will create new training set by concatenating 784 white noise dimensions to the existing 784 dimensions of the data
- we will also create an equivalent dataset by concatenating 784 all-zeros dimensions
- concatenation of meaningless features does not at all affect the information content of the data
- human classification accuracy wouldn't be affected by these transformations at all

```
from tensorflow.keras.datasets import mnist
import numpy as np

(train_images, train_labels), _ = mnist.load_data()
train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype("float32") / 255

train_images_with_noise_channels = np.concatenate( # add 784 features that are pure noise
    [train_images, np.random.random((len(train_images), 784))], axis=1)

train_images_with_zeros_channels = np.concatenate( # add 784 features that are all-zeros
    [train_images, np.zeros((len(train_images), 784))], axis=1)
```

Training the same model on MNIST data with noise channels or all-zero channels

```
from tensorflow import keras
from tensorflow.keras import layers

def get_model():
    model = keras.Sequential([
        layers.Dense(512, activation="relu"),
        layers.Dense(10, activation="softmax")
    ])
    model.compile(optimizer="rmsprop",
                  loss="sparse_categorical_crossentropy",
                  metrics=["accuracy"])
    return model

model = get_model()
history_noise = model.fit(
    train_images_with_noise_channels, train_labels,
    epochs=10,
    batch_size=128,
    validation_split=0.2) # use first 80% as training, rest 20% as validation

model = get_model()
history_zeros = model.fit(
    train_images_with_zeros_channels, train_labels,
    epochs=10,
    batch_size=128,
    validation_split=0.2)

Epoch 1/10
375/375 [=====] - 8s 19ms/step - loss: 0.6128 - accuracy: 0.8105 - val
Epoch 2/10
```

```

375/375 [=====] - 6s 17ms/step - loss: 0.2537 - accuracy: 0.9210 - val
Epoch 3/10
375/375 [=====] - 7s 19ms/step - loss: 0.1640 - accuracy: 0.9493 - val
Epoch 4/10
375/375 [=====] - 6s 17ms/step - loss: 0.1184 - accuracy: 0.9634 - val
Epoch 5/10
375/375 [=====] - 7s 19ms/step - loss: 0.0865 - accuracy: 0.9726 - val
Epoch 6/10
375/375 [=====] - 6s 16ms/step - loss: 0.0652 - accuracy: 0.9797 - val
Epoch 7/10
375/375 [=====] - 7s 19ms/step - loss: 0.0487 - accuracy: 0.9849 - val
Epoch 8/10
375/375 [=====] - 6s 16ms/step - loss: 0.0341 - accuracy: 0.9892 - val
Epoch 9/10
375/375 [=====] - 7s 19ms/step - loss: 0.0252 - accuracy: 0.9921 - val
Epoch 10/10
375/375 [=====] - 6s 16ms/step - loss: 0.0192 - accuracy: 0.9943 - val
Epoch 1/10
375/375 [=====] - 8s 20ms/step - loss: 0.2962 - accuracy: 0.9154 - val
Epoch 2/10
375/375 [=====] - 7s 18ms/step - loss: 0.1234 - accuracy: 0.9633 - val
Epoch 3/10
375/375 [=====] - 7s 19ms/step - loss: 0.0794 - accuracy: 0.9762 - val
Epoch 4/10
375/375 [=====] - 7s 20ms/step - loss: 0.0580 - accuracy: 0.9823 - val
Epoch 5/10
375/375 [=====] - 6s 16ms/step - loss: 0.0422 - accuracy: 0.9878 - val
Epoch 6/10
375/375 [=====] - 9s 23ms/step - loss: 0.0317 - accuracy: 0.9905 - val
Epoch 7/10
375/375 [=====] - 6s 16ms/step - loss: 0.0237 - accuracy: 0.9934 - val
Epoch 8/10
375/375 [=====] - 7s 19ms/step - loss: 0.0185 - accuracy: 0.9949 - val
Epoch 9/10
375/375 [=====] - 6s 16ms/step - loss: 0.0138 - accuracy: 0.9963 - val
Epoch 10/10
375/375 [=====] - 8s 23ms/step - loss: 0.0099 - accuracy: 0.9979 - val

```


Plotting a validation accuracy comparison

```

import matplotlib.pyplot as plt
val_acc_noise = history_noise.history["val_accuracy"]
val_acc_zeros = history_zeros.history["val_accuracy"]
epochs = range(1, 11)
plt.plot(epochs, val_acc_noise, "b-",
         label="Validation accuracy with noise channels")
plt.plot(epochs, val_acc_zeros, "b--",
         label="Validation accuracy with zeros channels")
plt.title("Effect of noise channels on validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

```

<matplotlib.legend.Legend at 0x7889c949ef50>



✓ The nature of generalization in deep learning

- deep learning models can be trained to fit anything, as long as they have enough representational power
- we will look at this "universal approximation" property later on in the course
- for now we will see a particular manifestation of this power
- 我们 shuffle the MNIST labels 然后基于这个乱七八糟的数据集来训练一个模型. 现在 inputs 和 shuffled labels 之间没有任何关系了.
- 即便如此, 并且即使模型相对较小, training loss 也会 decrease over time
- 但是自然地, the validation loss 不会随着 training loss 有任何改进.

✓ Fitting a MNIST model with randomly shuffled labels

```

(train_images, train_labels), _ = mnist.load_data()
train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype("float32") / 255

random_train_labels = train_labels[:]
np.random.shuffle(random_train_labels)

model = keras.Sequential([
    layers.Dense(512, activation="relu"),
    layers.Dense(10, activation="softmax")
])
model.compile(optimizer="rmsprop",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])
model.fit(train_images, random_train_labels,
          epochs=10,                                # increasing this will further increase the gap between
          batch_size=128,
          validation_split=0.2)

Epoch 1/10
375/375 [=====] - 10s 23ms/step - loss: 2.3166 - accuracy: 0.1044 - val
Epoch 2/10
375/375 [=====] - 8s 21ms/step - loss: 2.2993 - accuracy: 0.1174 - val
Epoch 3/10
375/375 [=====] - 5s 12ms/step - loss: 2.2921 - accuracy: 0.1269 - val
Epoch 4/10
375/375 [=====] - 7s 18ms/step - loss: 2.2801 - accuracy: 0.1386 - val
Epoch 5/10
375/375 [=====] - 4s 12ms/step - loss: 2.2648 - accuracy: 0.1501 - val
Epoch 6/10
375/375 [=====] - 7s 17ms/step - loss: 2.2448 - accuracy: 0.1653 - val
Epoch 7/10
375/375 [=====] - 8s 20ms/step - loss: 2.2217 - accuracy: 0.1811 - val
Epoch 8/10
375/375 [=====] - 7s 20ms/step - loss: 2.1946 - accuracy: 0.1983 - val
Epoch 9/10
375/375 [=====] - 5s 15ms/step - loss: 2.1643 - accuracy: 0.2136 - val
Epoch 10/10
375/375 [=====] - 5s 14ms/step - loss: 2.1312 - accuracy: 0.2321 - val
<keras.src.callbacks.History at 0x7889ca1ebf70>

```

✓ The manifold hypothesis(流形假设): deep learning 生效的根本原因


首先 manifold 的定义:(非严谨, 严谨定义参照 manifolds) A "manifold" is a lower-dimensional subspace of some parent space that is locally similar to a linear (Euclidean) space.

例如: a smooth curve in the plane is a 1D manifold within a 2D space, because for every point of the curve, you can draw a tangent. 又例如 a smooth surface within a 3D space is a 2D manifold

我们看到 MNIST 数据集.

- The space of 28 x 28 images with integer pixel values between 0 and 255 是一个比较大的 space, 而 MNIST samples 中的 handwritten digits 则是它的一个 tiny subspace.
- 这个 subspace 是 highly structured 的:
 - 它是 continuous 的: if you take a sample and modify it a little, it will still be recognizable as the same handwritten digit.

- 这个 **subspace** 中的所有 **samples** 都是 **connected by smooth paths that run through the subspace**. 即: 任取两个 random MNIST digits A and B, 都存在一系列的 “intermediate” images that morph A into B.
- 因而我们称: handwritten digits form a **manifold** within the space of possible 28×28 uint8 arrays



The manifold hypothesis 的具体内容: All natural data lies on a low-dimensional manifold within the high-dimensional space where it is encoded.

这是一个关于这个世界的信息结构的一个非常强的 statement. 并且我们认为这是 accurate 的, and it's the reason why deep learning works

- MNIST digits
- human faces
- tree morphology
- sounds of the human voice
- natural language

The manifold hypothesis implies:

- Machine learning models 只能把 relatively simple, low-dimensional, highly structured subspaces 给 fit 进它们的 potential input space (latent manifolds).
- 我们总是能够 *interpolate* between two inputs, that is to say, morph one into another via a continuous path along which all points fall on the manifold.
- The ability to interpolate between samples is the key to understanding generalization in deep learning.

begin side note

(*) it is natural to wonder what is the empirical evidence in favor of the manifold hypothesis. there is theoretical work on testing the manifold hypothesis:

[Testing the Manifold Hypothesis](#) by Charles Fefferman, Sanjoy Mitter, Hariharan Narayanan

but don't know if anyone has actually tested the hypothesis on empirical data. possible undergraduate thesis topic here?

end side note

Interpolation(插值) as a source of generalization

- if you have data points that can be interpolated, you can start making sense of points you've never seen before
- you can relate new points to other points that lie close on the manifold
- you can make sense of the totality of the space using only a sample of the space
- you can use interpolation to fill in the blanks

- interpolation on the latent manifold is different from linear interpolation in the parent space




Manifold interpolation
(intermediate point
on the latent manifold)




Linear interpolation
(average in the encoding space)

Why deep learning works: fit 和 overfit 的形象化


Before training:
the model starts
with a random initial state.




Beginning of training:
the model gradually
moves toward a better fit.




Further training: a robust
fit is achieved, transitively,
in the process of morphing
the model from its initial
state to its final state.




Final state: the model
overfits the training data,
reaching perfect training loss.



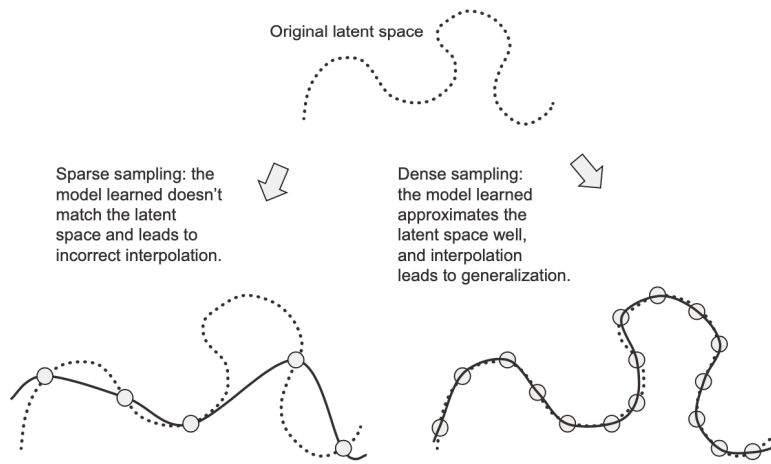
Test time: performance
of robustly fit model
on new data points



Test time: performance
of overfit model
on new data points



Training data is paramount



- suppose getting more data isn't possible
- the next best solution is to control model complexity

✓ Evaluating machine-learning models

- we will review regularization techniques soon

✓ Training, validation, and test sets

- evaluating a model always boils down to splitting the available data into three sets: training, validation, and test
- train on the training data and evaluate your model on the validation data
- once your model is ready for prime time, you test it one final time on the test data, which is meant to be as similar as possible to production data
- then you can deploy the model in production

Why not have two sets: a training set and a test set? You'd train on the training data and evaluate on the test data. Much simpler!

- developing a model always involves tuning its configuration:
 - choosing the number of layers
 - the size of the layers
- these are called *hyperparameters* to distinguish them from the parameters (which are the network's weights)
- this tuning is done by using as a feedback signal the performance of the model on the validation data

Simple hold-out validation

K-fold validation

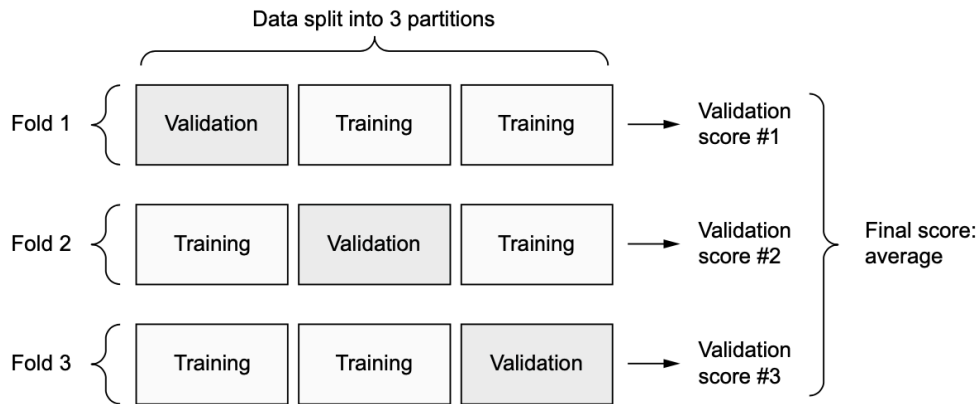


Figure 5.13 K-fold cross-validation with K=3

Iterated K-fold validation with shuffling

- this is for situations in which you have relatively little data available and you need to evaluate your model as precisely as possible
- consists of applying K-fold validation multiple times
- shuffle the data every time before splitting it K ways
- final score is the average of the scores obtained at each run of K-fold validation
- you end up training and evaluating $P * K$ models (where P is the number of iterations you use), which can be very expensive.

Beating a common-sense baseline

Important to ensure your trained network is doing better than some simple baseline

- in the MNIST digit-classification example, a simple baseline would be a validation accuracy greater than 0.1 (random classifier)
- in the IMDB example, it would be a validation accuracy greater than 0.5
- **class imbalance**: you have a binary classification problem where 90% of samples belong to class A and 10% belong to class B